

# WCET Analysis Lab: Assignment 1

Markus Klein  
Johannes Kasberger

SS 2012

## Problem 1

Recommended (3): As a warm-up exercise, follow the instruction in Timing Analysis Lab: First Steps

*Q: How long does it take to execute simple once, according to measurements, and according to the static analysis?*

## Answer

We used the compiler flag -O1.

- Measurements: 763 incl. function call - 70 overhead = 693 cycles
- Static analysis: 705 cycles

The difference is 12 cycles so the measurement is not far away from the static analysis. After one run we measured 1018 cycles but after reflashing the target the result was 763 cycles again. We can't explain this difference.

## Problem 2

Recommended (3): Also extract the instruction trace as outlined in Timing Analysis Lab: First Steps. Then compare the number of cycles needed in one iteration of the loop, with the number of cycles aiT calculated.

*Q: Do they coincide? What is the total number of cycles needed to execute simple according to the instruction trace buffer?*

## Answer

time	address	instruction	result
10839553	40001210	call 0x400011a0	[40001210]
10839568	40001214	st %g1, [%l1]	[40000000 07735939]
10839575	400011a0	mov 42, %g3	[0000002a]
10839582	400011a4	mov 0, %g2	[00000000]
10839589	400011a8	add %g3, 1, %g1	[0000002b]
10839596	400011ac	add %g1, 1, %g1	[0000002c]

10839603	400011b0	add	%g1, 1, %g1	[0000002d]
10839610	400011b4	add	%g1, 1, %g1	[0000002e]
10839617	400011b8	add	%g1, 1, %g1	[0000002f]
10839624	400011bc	add	%g1, 1, %g1	[00000030]
10839631	400011c0	add	%g1, 1, %g1	[00000031]
10839638	400011c4	add	%g1, 1, %g3	[00000032]
10839645	400011c8	add	%g2, 1, %g2	[00000001]
10839652	400011cc	cmp	%g2, 8	[ffffff9]
10839659	400011d0	bne	0x400011a8	[4000125c]
10839666	400011d4	nop		[00000000]
10839673	400011a8	add	%g3, 1, %g1	[00000033]
10839680	400011ac	add	%g1, 1, %g1	[00000034]
10839687	400011b0	add	%g1, 1, %g1	[00000035]
10839694	400011b4	add	%g1, 1, %g1	[00000036]
10839701	400011b8	add	%g1, 1, %g1	[00000037]
10839708	400011bc	add	%g1, 1, %g1	[00000038]
10839715	400011c0	add	%g1, 1, %g1	[00000039]
10839722	400011c4	add	%g1, 1, %g3	[0000003a]
10839729	400011c8	add	%g2, 1, %g2	[00000002]
10839736	400011cc	cmp	%g2, 8	[ffffffa]
10839743	400011d0	bne	0x400011a8	[4000125c]
10839750	400011d4	nop		[00000000]
10839757	400011a8	add	%g3, 1, %g1	[0000003b]
10839764	400011ac	add	%g1, 1, %g1	[0000003c]
10839771	400011b0	add	%g1, 1, %g1	[0000003d]
10839778	400011b4	add	%g1, 1, %g1	[0000003e]
10839785	400011b8	add	%g1, 1, %g1	[0000003f]
10839792	400011bc	add	%g1, 1, %g1	[00000040]
10839799	400011c0	add	%g1, 1, %g1	[00000041]
10839806	400011c4	add	%g1, 1, %g3	[00000042]
10839813	400011c8	add	%g2, 1, %g2	[00000003]
10839820	400011cc	cmp	%g2, 8	[ffffffb]
10839827	400011d0	bne	0x400011a8	[4000125c]
10839834	400011d4	nop		[00000000]
10839841	400011a8	add	%g3, 1, %g1	[00000043]
10839848	400011ac	add	%g1, 1, %g1	[00000044]
10839855	400011b0	add	%g1, 1, %g1	[00000045]
10839862	400011b4	add	%g1, 1, %g1	[00000046]
10839869	400011b8	add	%g1, 1, %g1	[00000047]
10839876	400011bc	add	%g1, 1, %g1	[00000048]
10839883	400011c0	add	%g1, 1, %g1	[00000049]
10839890	400011c4	add	%g1, 1, %g3	[0000004a]
10839897	400011c8	add	%g2, 1, %g2	[00000004]
10839904	400011cc	cmp	%g2, 8	[ffffffc]
10839911	400011d0	bne	0x400011a8	[4000125c]
10839918	400011d4	nop		[00000000]

10839925	400011a8	add	%g3, 1, %g1	[0000004b]
10839932	400011ac	add	%g1, 1, %g1	[0000004c]
10839939	400011b0	add	%g1, 1, %g1	[0000004d]
10839946	400011b4	add	%g1, 1, %g1	[0000004e]
10839953	400011b8	add	%g1, 1, %g1	[0000004f]
10839960	400011bc	add	%g1, 1, %g1	[00000050]
10839967	400011c0	add	%g1, 1, %g1	[00000051]
10839974	400011c4	add	%g1, 1, %g3	[00000052]
10839981	400011c8	add	%g2, 1, %g2	[00000005]
10839988	400011cc	cmp	%g2, 8	[fffffffd]
10839995	400011d0	bne	0x400011a8	[4000125c]
10840002	400011d4	nop		[00000000]
10840009	400011a8	add	%g3, 1, %g1	[00000053]
10840016	400011ac	add	%g1, 1, %g1	[00000054]
10840023	400011b0	add	%g1, 1, %g1	[00000055]
10840030	400011b4	add	%g1, 1, %g1	[00000056]
10840037	400011b8	add	%g1, 1, %g1	[00000057]
10840044	400011bc	add	%g1, 1, %g1	[00000058]
10840051	400011c0	add	%g1, 1, %g1	[00000059]
10840058	400011c4	add	%g1, 1, %g3	[0000005a]
10840065	400011c8	add	%g2, 1, %g2	[00000006]
10840072	400011cc	cmp	%g2, 8	[fffffffe]
10840079	400011d0	bne	0x400011a8	[4000125c]
10840086	400011d4	nop		[00000000]
10840093	400011a8	add	%g3, 1, %g1	[0000005b]
10840100	400011ac	add	%g1, 1, %g1	[0000005c]
10840107	400011b0	add	%g1, 1, %g1	[0000005d]
10840114	400011b4	add	%g1, 1, %g1	[0000005e]
10840121	400011b8	add	%g1, 1, %g1	[0000005f]
10840128	400011bc	add	%g1, 1, %g1	[00000060]
10840135	400011c0	add	%g1, 1, %g1	[00000061]
10840142	400011c4	add	%g1, 1, %g3	[00000062]
10840149	400011c8	add	%g2, 1, %g2	[00000007]
10840156	400011cc	cmp	%g2, 8	[fffffff]
10840163	400011d0	bne	0x400011a8	[4000125c]
10840170	400011d4	nop		[00000000]
10840177	400011a8	add	%g3, 1, %g1	[00000063]
10840184	400011ac	add	%g1, 1, %g1	[00000064]
10840191	400011b0	add	%g1, 1, %g1	[00000065]
10840198	400011b4	add	%g1, 1, %g1	[00000066]
10840205	400011b8	add	%g1, 1, %g1	[00000067]
10840212	400011bc	add	%g1, 1, %g1	[00000068]
10840219	400011c0	add	%g1, 1, %g1	[00000069]
10840226	400011c4	add	%g1, 1, %g3	[0000006a]
10840233	400011c8	add	%g2, 1, %g2	[00000008]
10840240	400011cc	cmp	%g2, 8	[00000000]

10840247	400011d0	bne	0x400011a8	[4000125c]
10840254	400011d4	nop		[00000000]
10840263	400011d8	retl		[400011d8]

Difference: 710 cycles

Just based on the number of instructions we calculated the number of cycles needed = 118 cycles. This is much shorter than the result of the static analysis. The reason for this is that we don't know the exact memory timing and the instruction decode takes more time.

### Problem 3

Mandatory (4): First, create a project containing the files contained in the insertion sort folder of the task specification. Now complete the function `main.c:run()`, executing insertion sort a few times, with array size 32 and different input data. Measure the minimum and maximum time needed to execute the sort function.

*Q: What were the results of the measurement? How many test sets would do you need to cover all possible execution path?*

### Answer

We used the compiler flag `-Os`.

Case	measurement[cycles]
Best-case (pre sorted)	3659
Worst-case (upside down sorted)	44267
Average-case (unsorted)	24875

To cover all possible execution paths to sort a array of 32 32 bit integers we need maximal  $2^{32} * 32 = 137438953472$  test cases.

### Problem 4

Mandatory (5): Add loop bounds and additional flow facts for `insertion sort.c:insertion sort()`, using the symbolic name `@size` for the size of the array to be sorted. Next, analyze the WCET of insertion sort, assuming an array size of 32. Keep the array size as a symbolic name (user register `@size`). Finally, write a test function which calls insertion sort more than once, with different array sizes (e.g., 16,32 and 64). Also repeat the static analysis with different array sizes.

*Q: How many cycles do you need to execute insertion sort according to the static analysis?*

### Answer

The static analysis resulted in 74694 cycles.

*Q: What results do you get for an array size of 16, 32 or 64, using measurements and static analysis?*

**Answer**

method	size 16[cycles]	size 32[cycles]	size 64[cycles]
measurement	11527	34579	132899
static	18232	77224	318088

*Q: In addition to the size of the array, what other aspects of the input data might influence the WCET?*

**Answer**

The structure of the data. For instance: The worst-case occurs when the input data is sorted upside down. The best-case occurs if the input data is already sorted.

### Problem 5

Recommended (8 Points): Assume that your goal was to find out the WCET of `task.c:task()`. Before analyzing the execution time, you should answer a few basic questions about the input data for the monitoring task, and analyze the control flow on the source code level.

*Q: What is the set of input data which might influence the execution time of the task at the software side?*

**Answer**

If many samples are missing the interpolation of the data has big influence on the runtime. Most of the other loops just depend on the amount of samples used for calculation but this does not depend on the input data and is determined at compile time.

*Q: Is it tractable to enumerate every possible input?*

**Answer**

No it's not traceable. The result of the interpolation is dependent on the history of the last samples so enumerating the history and the inputs would require many different testcases.

*Q: Which loops need to be bounded?*

**Answer**

The loops we found are in the merge samples function.

*Q: Add all loop bounds and flow facts you can find to the file `task.c` (as source code annotations).*

## Answer

```
for(i = valid+1;
    i < cnt; /* ai: loop here MAX (@inputcount+4); */
    i++)
{
    if(i >= 0)
    {
        x = xs[i];
        /* ai: flow (here) <= 32 ("merge_samples"); */
        sample_buffer_set(sbuf, i, x);
    }
    else
    {
        x = sample_buffer_get(sbuf, i);
        /* ai: flow (here) <= 4 ("merge_samples"); */
    }

    /* If the sample is not missing, interpolate the ones before if the range is acceptable */
    if(! IS_VALUE_MISSING(x))
    {
        /* Only interpolate if we interpolate at most MAX_CONSECUTIVE_MISSING samples */
        int missing_samples = i - valid - 1;
        if(missing_samples > 0 && missing_samples <= MAX_CONSECUTIVE_MISSING)
        {
            for(j = i-1; j > valid; --j)
            {
                /* At most once for each invalid input sample */

                sample_value_t y = sample_buffer_get(sbuf, j);
                /* ai: loop here MAX 4; */
                /* ai: flow each (here) / ("merge_samples") is max 32; */
                if(! IS_VALUE_MISSING(y)) break;
                y = iinterpolate16(valid, sample_buffer_get(sbuf, valid), i, x, j);
                sample_buffer_set(sbuf, j, y);
            }
            valid = i;
        }
    }
}
```

## Problem 6

Recommended (8 Points): Analyze the `fft()` function called in `task.c`. Try to find loop bounds for the Fast Fourier Transform implementation (`fixedpoint.c:fp_radix2fft` with-scaling) first. If you have difficulties finding them, add a debug statement and run the transform with different input data sizes. Add flow constraints relating the execution frequency of the inner loops with the functions execution frequency. Finally, try to analyze the execution time using `aiT`. There is already a timing measurement for the `fft` in

the executable, so it is easy to compare the number of cycles estimated to execute the function.

*Q: Compare the worst-case number of iterations for the inner loop with and without using these flow constraints. Finally, think about the complexity of calculating loop bounds for FFT.*

### Answer

Comparison static analysis with measurement:

type	cycles
measurement	178807 - 124755
static with flow constraints	160344 - 159957

Comparison of inner loop in fp radix2fft withscaling with/without flow constraints:

type	cycles
static without flow constraints	1998624
static with flow constraints	79491

Without the flow constraint the inner loops have a big impact on WCET. The reason is that each inner loop is at most executed 32 times in each iteration but this number changes with every iteration. We calculated the flow constraint with  $\sum_{i=0}^{t-1} 2^i * 2^{t-i} = 192$  to represent this fact.

*Q: Does the FFT loop bound depend on the input data?*

### Answer

It just depends on the amount of the input data and not the data itself.

### Problem 7

Optional Challenge (5 Bonus Points): Try to analyze the WCET of task.c:task() using aiT. If you attempt to solve this challenge, use the control flow graph and disassembling capabilities of aiT, and be sure to understand the source code you are analyzing.

### Problem 8

Mandatory (4): Answer the following questions

*Q: How much time did you spend writing annotations and analyzing the code? Was it less or more than you expected? How much time did you spend on this first assignment?*

### Answer

- Mandatory Part: 6 hours
- Recommended Part: 6 hours
- Protocol: 2 hours

We worked 14 hours on this assignment. It was more than expected, because we ran into many problems which were not directly related to the given exercises (e.g. configuration errors in aiT and so on).

*Q: What is the ratio between observed and calculated execution time? Discuss the causes of the overestimation.*

**Answer**

For the static analysis we encountered a result that sometime was twice as long as our measurements. The causes of overestimation are too less restrictive loop bounds and flow annotations. The memory access times of the CPU model are a pessimistic estimation and can also be a source of overestimation.

*Q: As you learned, sometimes it is necessary to annotate the assembler code. Why? What problems can you see because of this?*

**Answer**

The code has another structure in assembler after the compiler has finished its optimisation. The source code annotation points to a line of code that is on another position in the assembler code. Sometimes the loop body is moved before the loop condition and then it's also hard to see where a source code annotation belongs to.

The problem of annotating the assembler code is that every little change and recompilation of a program can possibly result in totally different binaries and would require to adapt the assembler annotations to this new binary.