

WCET Analysis Lab: Assignment 3

Markus Klein
Johannes Kasberger

SS 2012

Problem 1

Recommended, 8 Points: As a first step, implement a test driver, which calls the chosen benchmark with different input data. Compile the test driver (using -Os), and measure the execution time. Next, add loop bounds and other flow facts to the selected benchmark, and analyze the WCET using aiT.

Q: Measurement Result

Answer

Our measurement is based on a test case with input length 26 bytes.

Original benchmark output:

Measurement: Due to prescaling, we have cycles = 4 * ticks +- 3

Assumed Overhead : 4 (max:14, cold: 18)

Assumed Overhead (with function call): 26 (max:30, cold: 70)

Sig correct

md5 estimated: 9135 [incl.function call]

The runtime of this benchmark was 9135 cycles in contrast to the WCET calculated by aiT in a static analysis of 9266 cycles.

Problem 2

Recommended, 8 Points: Next, apply the single-path transformation to this implementation. In theory, loop bounds are unnecessary for single path code, and, on a suitable platform, execution time variations should be minimal.

Q: Investigate whether this is the case, using the test driver developed in the previous problem, and using aiT.

Answer

After the transformation of the code into single-path form the same benchmark had a

runtime of 25323 cycles, which is nearly three times the original runtime. The static analysis yields a WCET of 42944 cycles.

To the best of our knowledge we're not able to find a reason for this huge discrepancy between the actual runtime and the WCET. All (automatically determined) loop bounds seem to be right in aiT.

As we can see in Figure 1 there is quite a difference in the steepness of the runtime curve, when testing with input length 1 up to 56. The raw data can be found in Table 1.

The variation of the execution time is 92 cycles compared to the original code, which has a variation of 2380 cycles.

Problem 3

Recommended, 11 Points: Next, implement a WCET-oriented solution for the given problem, which has a good worst-case performance, automatically analyzable control flow, and, assuming a suitable processor, a small execution time variance. To this end, it is acceptable to add additional parameters to functions, and to use different algorithms.

Q: Again, perform measurements and a WCET analysis, and summarize all collected results in a table in your report.

Answer

After restricting the possible input size of the MD5 hash function to a maximum of 56 bytes, we could remove big parts of code and therefore further drill down the execution time.

The benchmark took 11883 cycles and aiT predicts a WCET of 17778 cycles.

As Figure 1 shows, the optimized version behaves similar to the single-path version when run over the bigger test set, but has roughly half of the WCET. The raw data can be found in Table 1.

The variation of the execution time is 224 cycles compared to single-path code, which has a variation of 92 cycles.

Problem 4

Recommended, 8 Points: Finally, specify the WCET of your library function (third version) as precisely as possible. Your deliverable should include the object code of the WCET-oriented solution, along with a header file which contains the function prototype, as well as your test driver and a documented .ais file for the function. Provide detailed explanations how calls to your function will influence the execution time of the program it is embedded in, how to estimate the functions execution time using measurements, and how to analyze it using aiT. We will evaluate your function (using measurements and aiT), without inspecting the code and without adding any new flow facts. Furthermore, note that in this assignment, LEON3 is synthesized with an LRU instruction cache; therefore, the execution time may vary depending on the cache state.

Q: Description

Answer

Object file: md5.o

Header file: md5.h

Testdriver: measure.c

WCET: 17778 cycles

aiT: No annotations needed.

Measured runtimes for all possible input lengths can be seen in Table 1.

In our case the function call_md5 was mapped to memory address 0x40001218. Depending on the exact location of the function the measured results may vary as the instruction cache may experience different cache line allocations.

The runtime of the routine has only small variations.

The bound calculated by the static analysis will hold for any placement, as the calculation was started with unknown cache state.

Calls to our function will manipulate the program it is embedded in mainly in the instruction cache behavior.

Since the program is in single-path form it can be analyzed by just running aiT using the given project file or targeting the function call_md5 in measure.c.




Figure 1: Comparison between the three versions of the program

Table 1: Raw measurements for big test

Length	Original	Single Path	Optimized
1	8875	25219	11799
2	8887	25219	11799
3	8895	25223	11807
4	8907	25223	11811
5	8915	25223	11811
6	8923	25227	11815
7	8935	25223	11819
8	8947	25227	11819
9	8955	25227	11823
10	8963	25231	11827
11	8975	25227	11827
12	8987	25231	11831
13	8995	25231	11835
14	9007	25231	11839
15	9015	25231	11839
16	9027	25235	11847
17	9039	25235	11847
18	9043	25235	11851
19	9059	25239	11851
20	9067	25239	11855
21	9079	25239	11859
22	9087	25239	11863
23	9095	25243	11863
24	9107	25239	11871
25	9119	25247	11871
26	9127	25247	11871
27	9135	25243	11879
28	9147	25243	11879
29	9155	25247	11883
30	9167	25247	11887
31	9179	25251	11891
32	9183	25247	11895
33	9195	25251	11895
34	9207	25255	11895
35	9215	25255	11903
36	9227	25251	11903
37	9235	25255	11907
38	9247	25255	11911
39	9259	25255	11915
40	9267	25259	11915
41	9275	25259	11919
42	9287	25259	11919
43	9295	25263	11927

Table 1: Raw measurements for big test

Length	Original	Single Path	Optimized
44	9303	25263	11927
45	9315	25263	11931
46	9327	25263	11935
47	9339	25267	11939
48	9343	25267	11939
49	9359	25267	11943
50	9367	25271	11943
51	9375	25271	11947
52	9387	25271	11951
53	9395	25271	11955
54	9403	25275	11959
55	9351	25267	11955
56	11255	25311	12023
Difference Max-Min	2380	92	224